

CSE 6140/ CX 4140

Computational Science and Engineering

ALGORITHMS

Greedy Algorithms - 4


Instructor: Xiuwei Zhang

School of Computational Science and Engineering

Based on slides by Prof. Ümit V. Çatalyürek

Dijkstra's algorithm: proof of correctness (greedy stays ahead)

- **Invariant.** For each node $u \in S$: $d[u]$ = length of a shortest $s \rightsquigarrow u$ path.
- **Pf.** [by induction on $|S|$]
- **Base case:** $|S| = 1$ is easy since $S = \{ s \}$ and $d[s] = 0$.
- **Inductive hypothesis:** Assume true for $|S| \geq 1$.
 - Let v be next node added to S , and let (u, v) be the final edge.
 - A shortest $s \rightsquigarrow u$ path plus (u, v) is an $s \rightsquigarrow v$ path of length $\pi(v)$.
 - Consider **any** other $s \rightsquigarrow v$ path P . We show that it is no shorter than $\pi(v)$.
 - Let $e = (x, y)$ be the first edge in P that leaves S , and let P' be the subpath from s to x .
 - The length of P is already $\geq \pi(v)$ as soon as it reaches y :



$$\ell(P) \geq \ell(P') + \ell_e \geq d[x] + \ell_e \geq \pi(y) \geq \pi(v) \quad \blacksquare$$

\uparrow
 non-negative
lengths

\uparrow
 inductive
hypothesis

\uparrow
 definition
of $\pi(y)$

\uparrow
 Dijkstra chose v
instead of y


Dijkstra's algorithm demo

- Initialize $S \leftarrow \{ s \}$ and $d[s] \leftarrow 0$.
- Repeatedly choose unexplored node $v \notin S$ which minimizes

$$\pi(v) = \min_{e=(u,v) : u \in S} d[u] + \ell_e$$

the length of a shortest path from s to some node u in explored part S , followed by a single edge $e = (u, v)$

add v to S ; set $d[v] \leftarrow \pi(v)$ and $pred[v] \leftarrow \operatorname{argmin}$.



Redundant
calculation?


Dijkstra's algorithm demo

- Initialize $S \leftarrow \{ s \}$ and $d[s] \leftarrow 0$.
- Repeatedly choose unexplored node $v \notin S$ which minimizes

$$\pi(v) = \min_{e=(u,v) : u \in S} d[u] + \ell_e$$

the length of a shortest path from s to some node u in explored part S , followed by a single edge $e = (u, v)$

add v to S ; set $d[v] \leftarrow \pi(v)$ and $pred[v] \leftarrow \operatorname{argmin}$.




Dijkstra's algorithm demo

- Initialize $S \leftarrow \{ s \}$ and $d[s] \leftarrow 0$.
- Repeatedly choose unexplored node $v \notin S$ which minimizes

$$\pi(v) = \min_{e=(u,v) : u \in S} d[u] + \ell_e$$

the length of a shortest path from s to some node u in explored part S , followed by a single edge $e = (u, v)$

add v to S ; set $d[v] \leftarrow \pi(v)$ and $pred[v] \leftarrow \operatorname{argmin}$.



Redundant
calculation?


Dijkstra's algorithm demo

- Initialize $S \leftarrow \{ s \}$ and $d[s] \leftarrow 0$.
- Repeatedly choose unexplored node $v \notin S$ which minimizes

$$\pi(v) = \min_{e=(u,v) : u \in S} d[u] + \ell_e$$

the length of a shortest path from s to some node u in explored part S , followed by a single edge $e = (u, v)$

add v to S ; set $d[v] \leftarrow \pi(v)$ and $pred[v] \leftarrow \operatorname{argmin}$.




Dijkstra's algorithm: efficient implementation

- **Critical optimization 1.** For each node $v \notin S$: explicitly maintain $\pi[v]$ instead of computing directly from definition

$$\pi(v) = \min_{e = (u,v) : u \in S} d[u] + \ell_e$$

- For each $v \notin S$: $\pi(v)$ can only decrease (because set S increases).
- More specifically, suppose u is added to S and there is an edge $e = (u, v)$ leaving u . Then, it suffices to update:

$$\pi[v] \leftarrow \min \{ \pi[v], \pi[u] + \ell_e \}$$

 recall: for each $u \in S$,
 $\pi[u] = d[u] = \text{length of shortest } s \rightsquigarrow u \text{ path}$

- **Critical optimization 2.** Use a min-oriented **priority queue** (PQ) to choose an unexplored node that minimizes $\pi[v]$.

Dijkstra's algorithm: efficient implementation

- **Implementation.**
 - Algorithm maintains $\pi[v]$ for each node v .
 - Priority queue stores unexplored nodes, using $\pi[\cdot]$ as priorities.
 - Once u is deleted from the PQ, $\pi[u]$ = length of a shortest $s \rightsquigarrow u$ path.

DIJKSTRA (V, E, ℓ, s)

FOREACH $v \neq s : \pi[v] \leftarrow \infty, \text{pred}[v] \leftarrow \text{null}; \pi[s] \leftarrow 0.$

Create an empty priority queue PQ .

FOREACH $v \in V : \text{INSERT}(PQ, v, \pi[v]).$

WHILE (*IS-NOT-EMPTY*(PQ))

$u \leftarrow \text{DEL-MIN}(PQ).$

FOREACH edge $e = (u, v) \in E$ leaving u :


IF ($\pi[v] > \pi[u] + \ell_e$)

$\pi[v] \leftarrow \pi[u] + \ell_e; \text{pred}[v] \leftarrow e.$

DECREASE-KEY($PQ, v, \pi[v]$).

Dijkstra's Shortest Path Algorithm

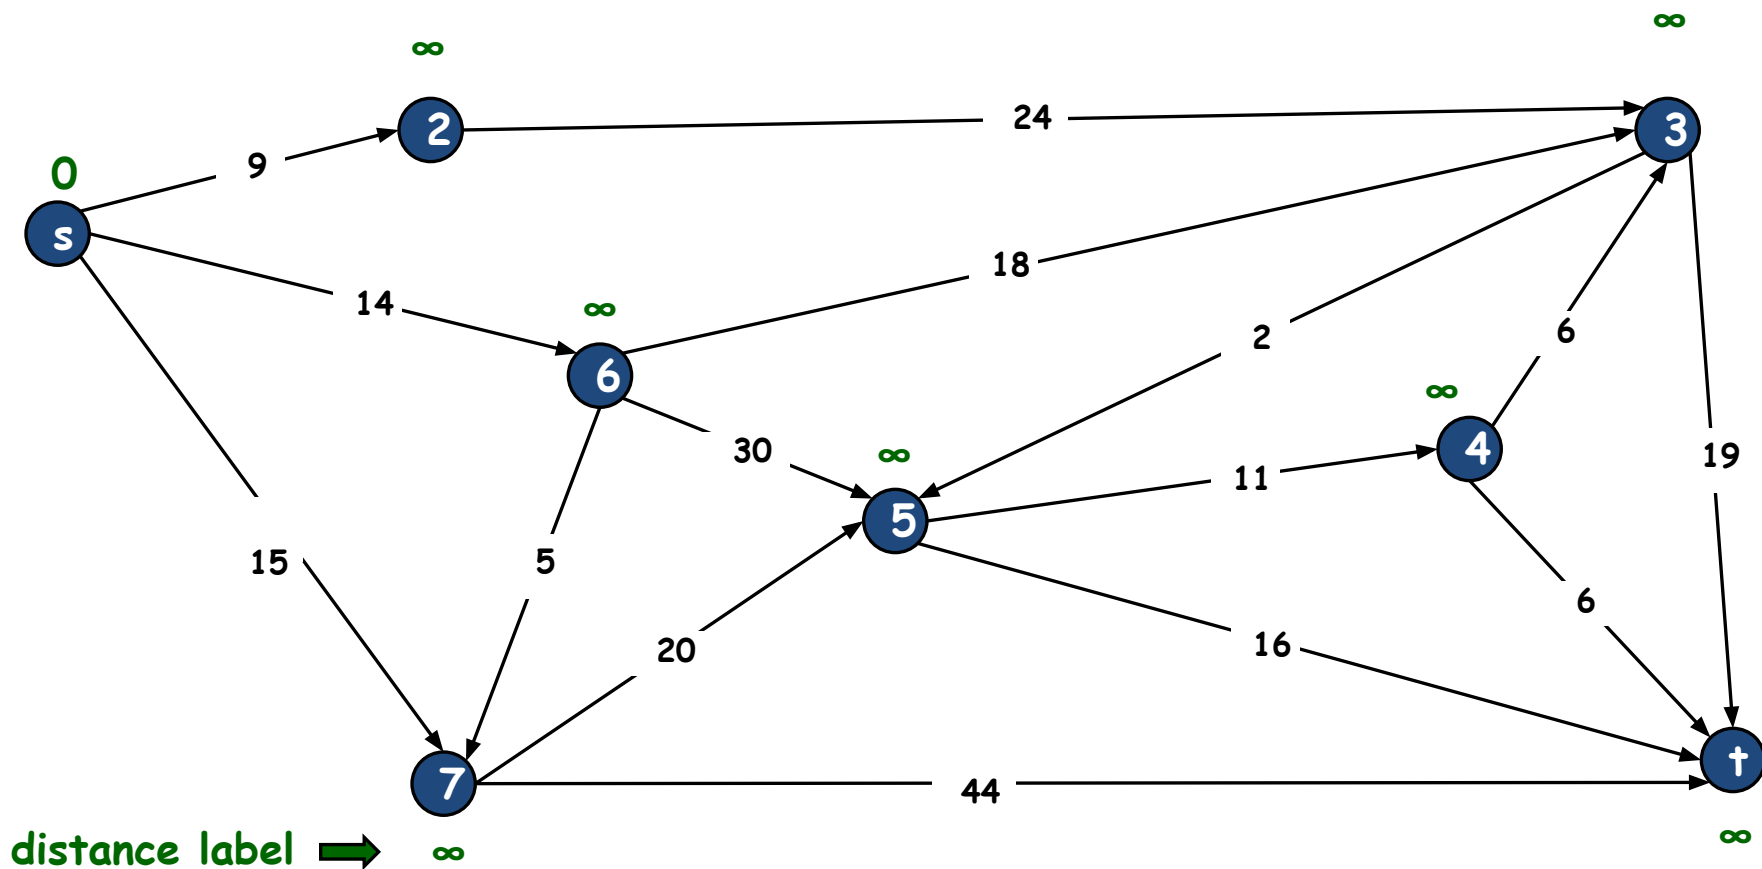
Find shortest path from s to t.



Dijkstra's Shortest Path Algorithm

$S = \{ \}$

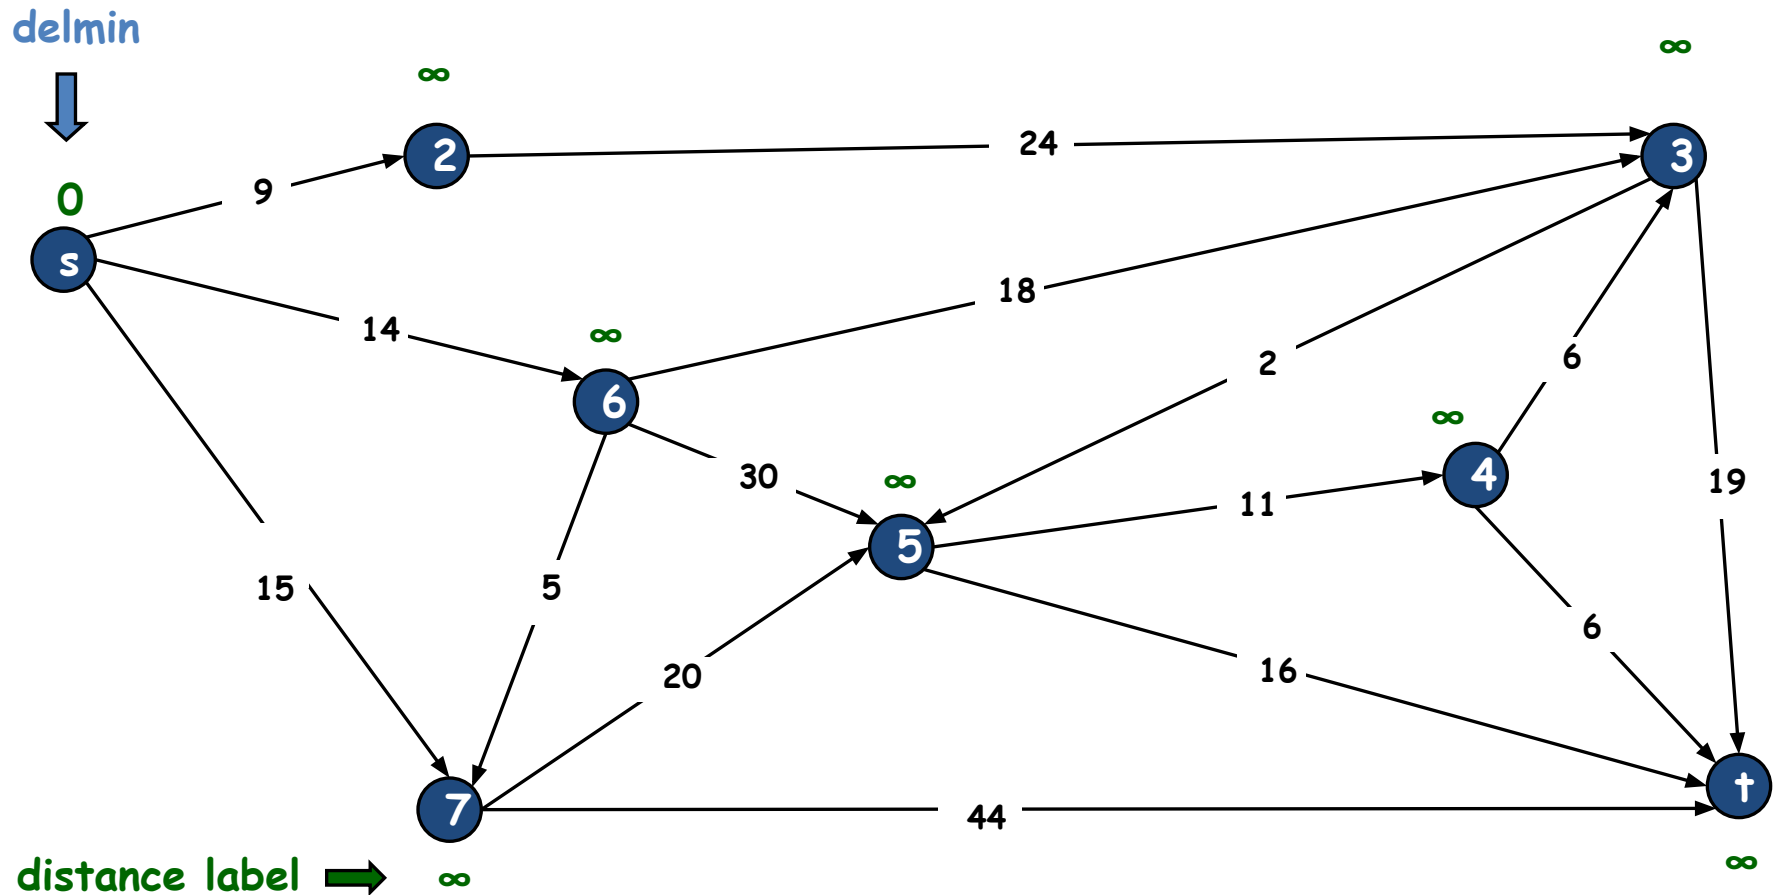
$PQ = \{ s, 2, 3, 4, 5, 6, 7, t \}$



Dijkstra's Shortest Path Algorithm

$S = \{ \}$

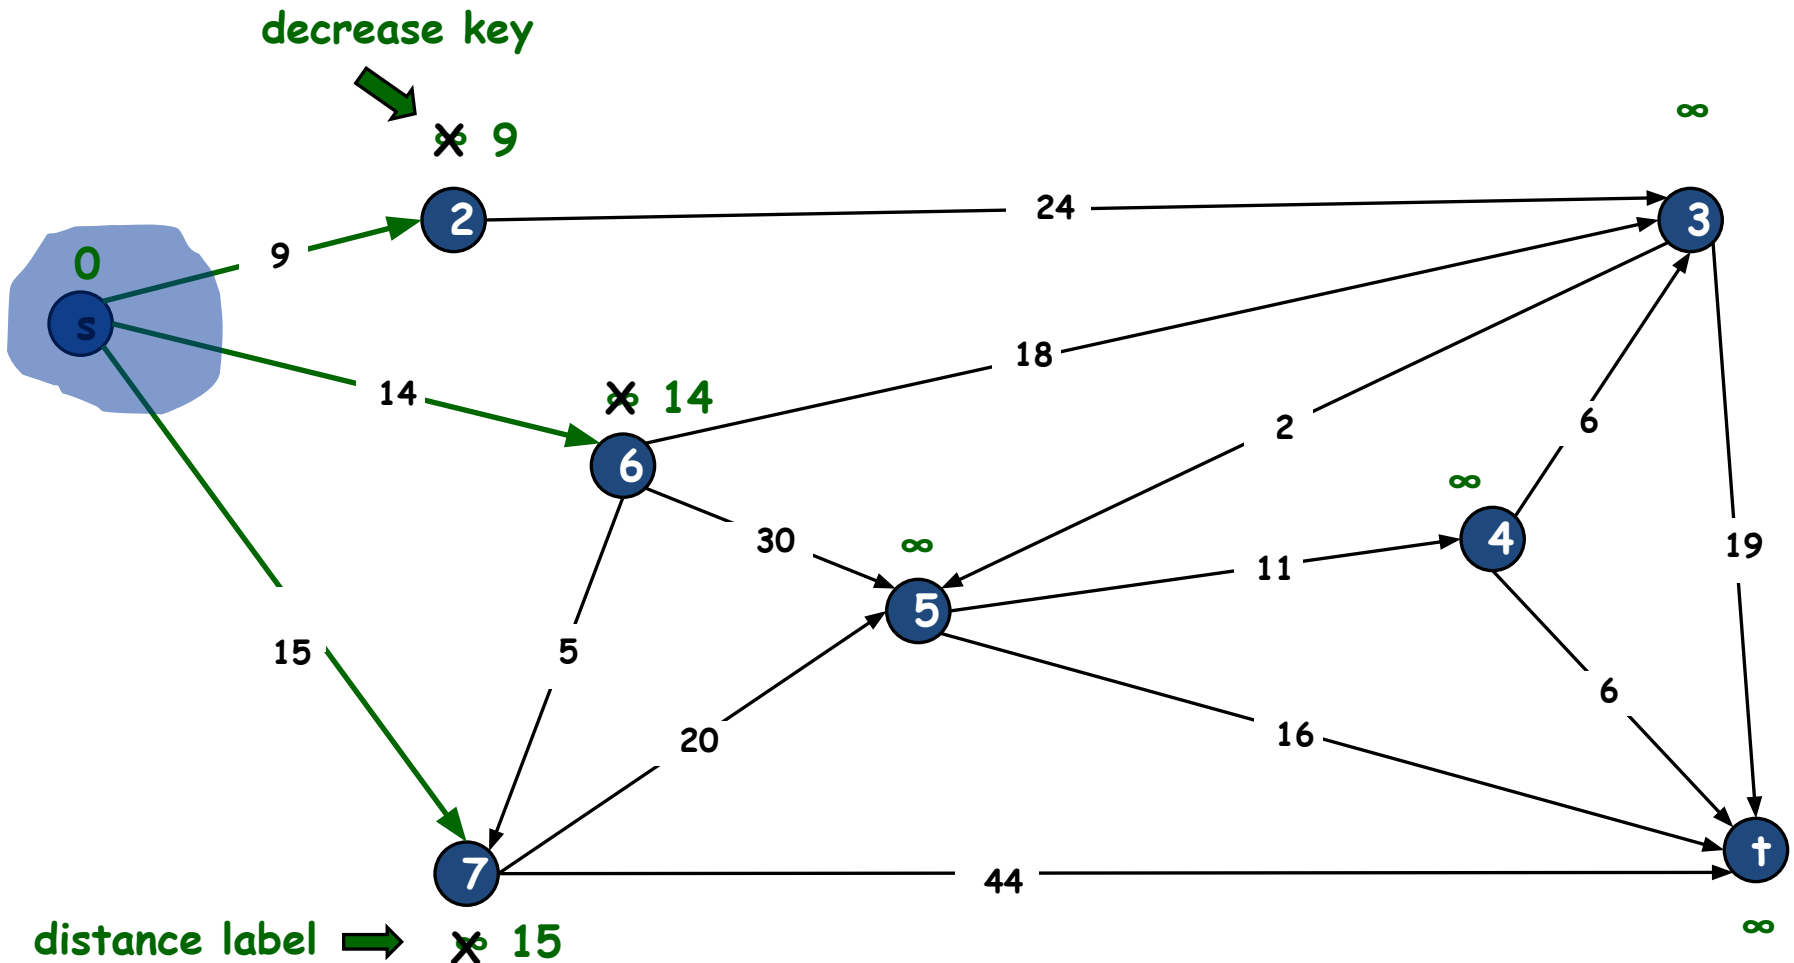
$PQ = \{ s, 2, 3, 4, 5, 6, 7, t \}$



Dijkstra's Shortest Path Algorithm

$S = \{ s \}$


$PQ = \{ 2, 3, 4, 5, 6, 7, \dagger \}$



Dijkstra's Shortest Path Algorithm

$S = \{s\}$

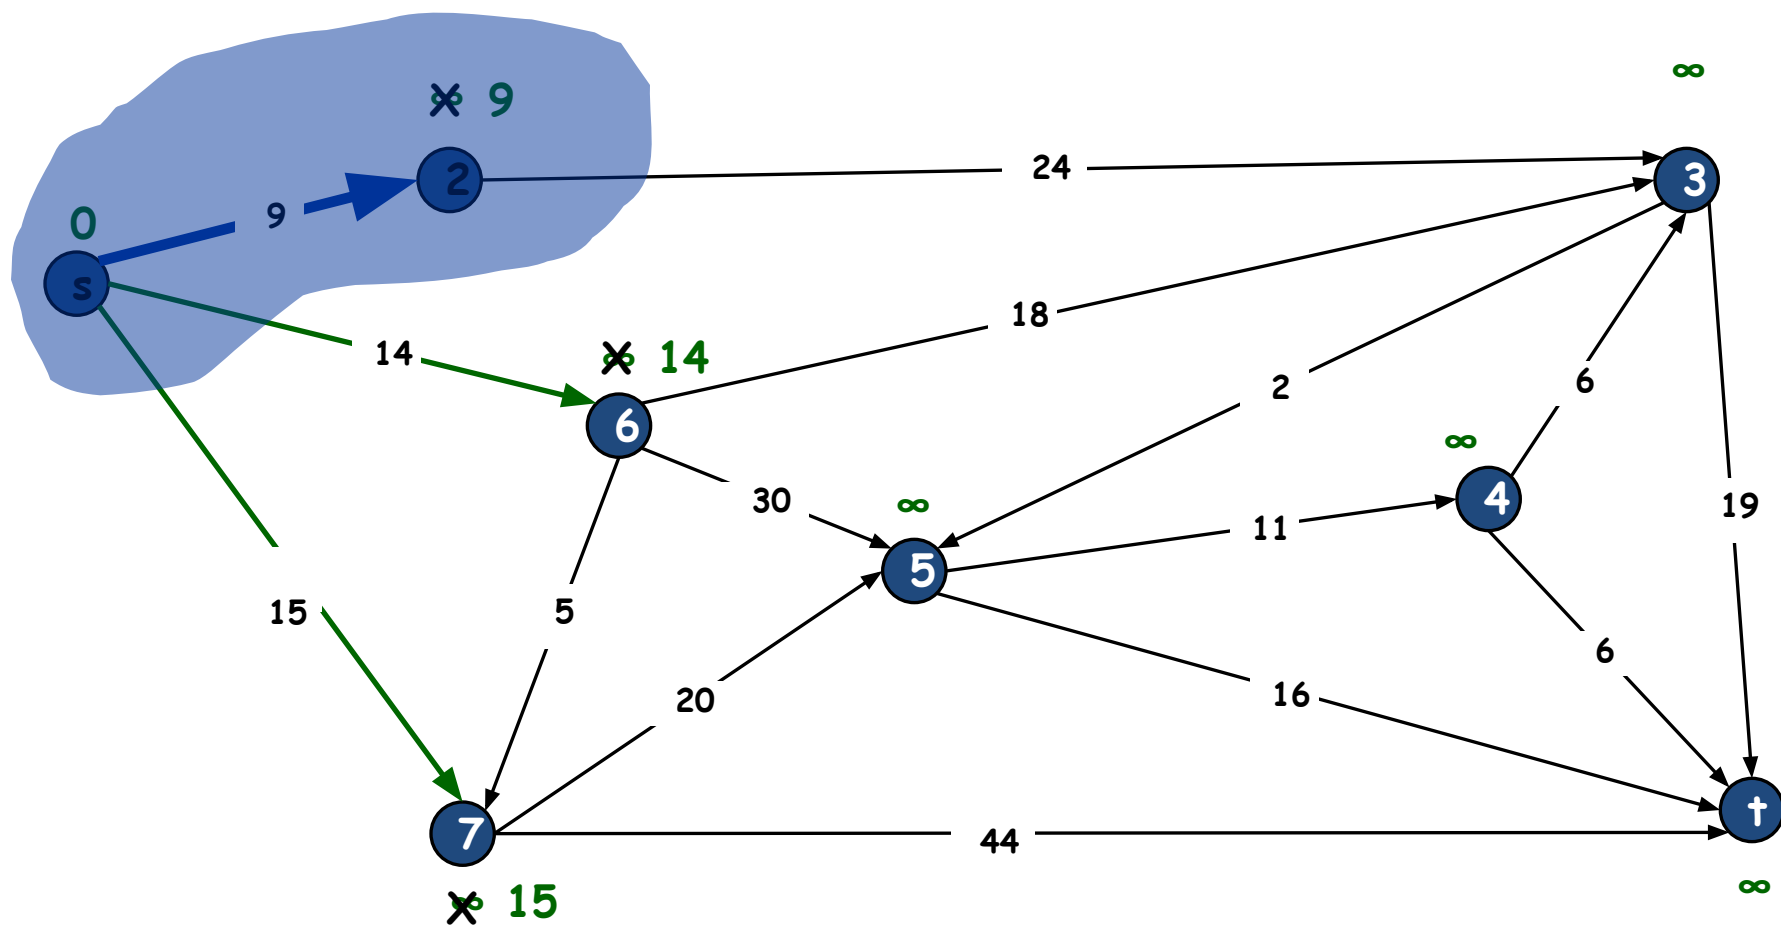
$PQ = \{2, 3, 4, 5, 6, 7, \dagger\}$



Dijkstra's Shortest Path Algorithm

$S = \{s, 2\}$

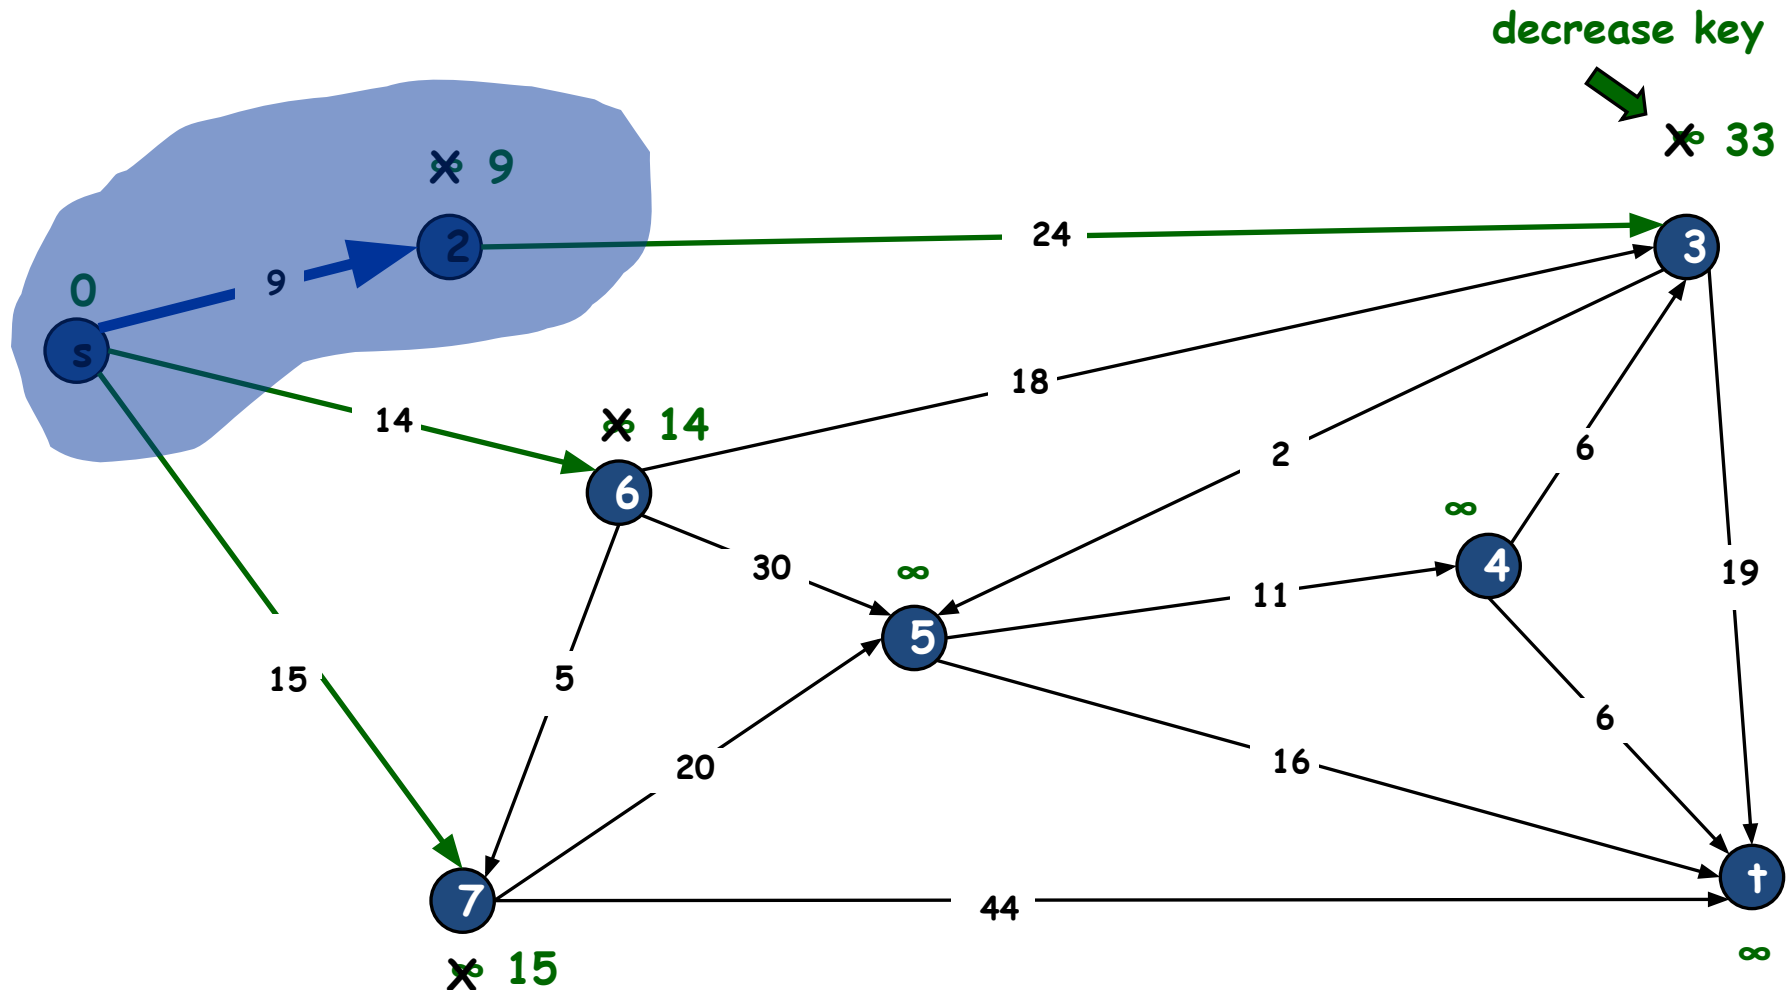
$PQ = \{3, 4, 5, 6, 7, \dagger\}$



Dijkstra's Shortest Path Algorithm

$S = \{s, 2\}$

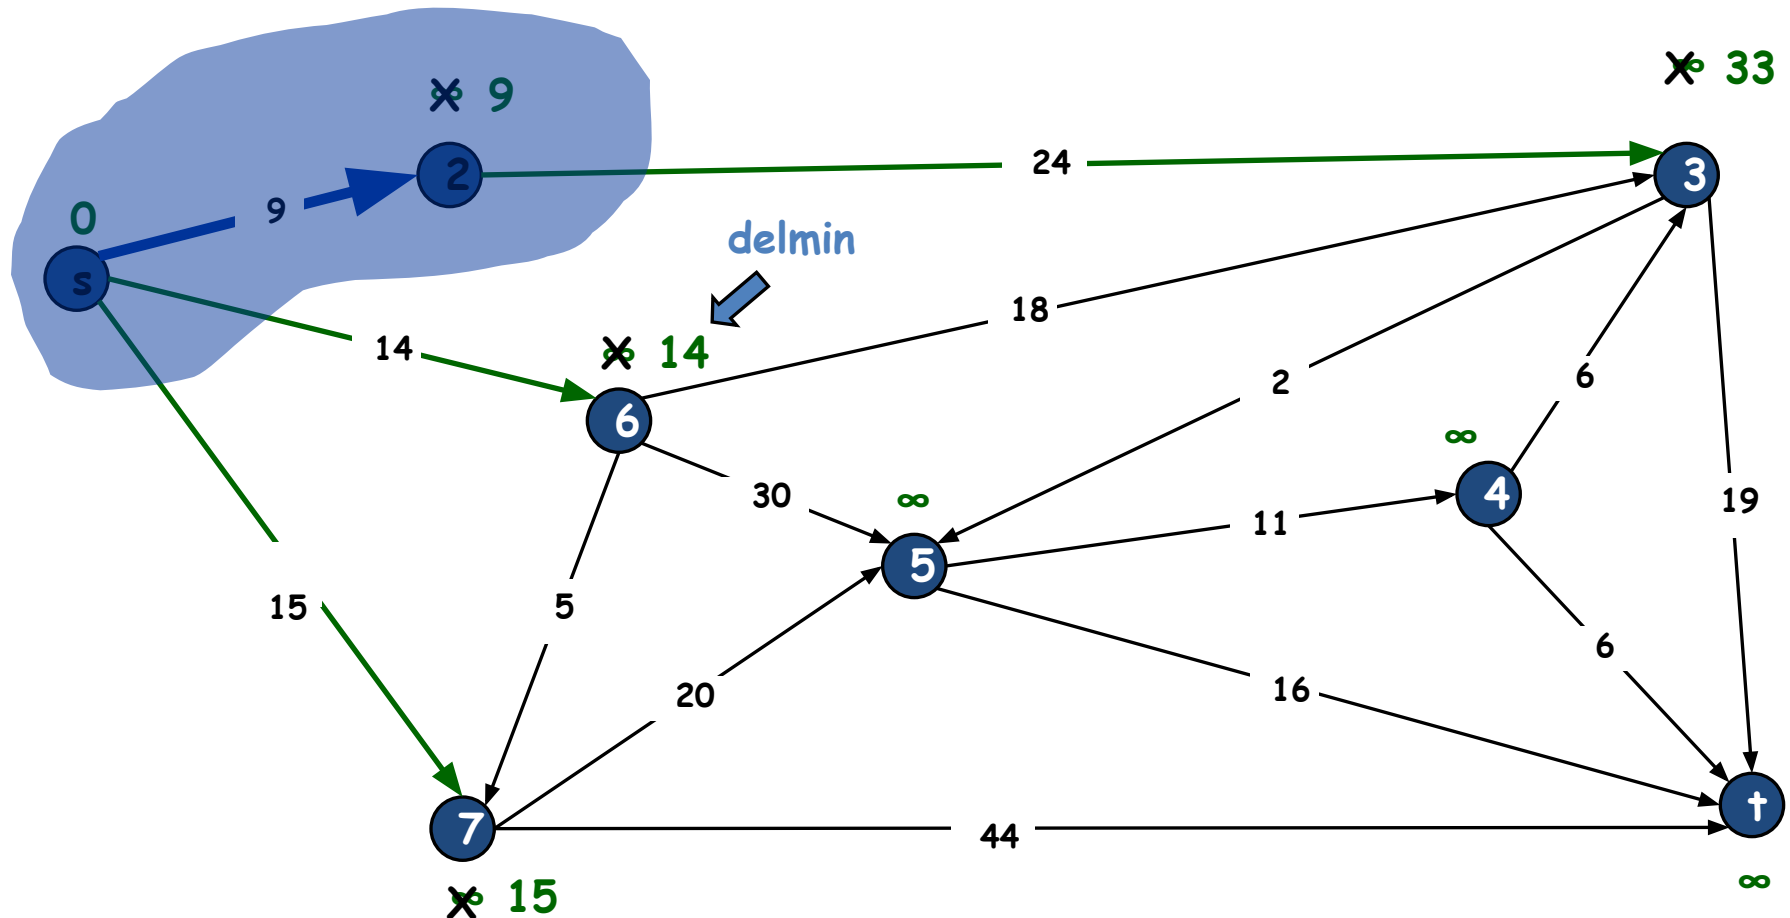
$PQ = \{3, 4, 5, 6, 7, \dagger\}$



Dijkstra's Shortest Path Algorithm

$S = \{s, 2\}$

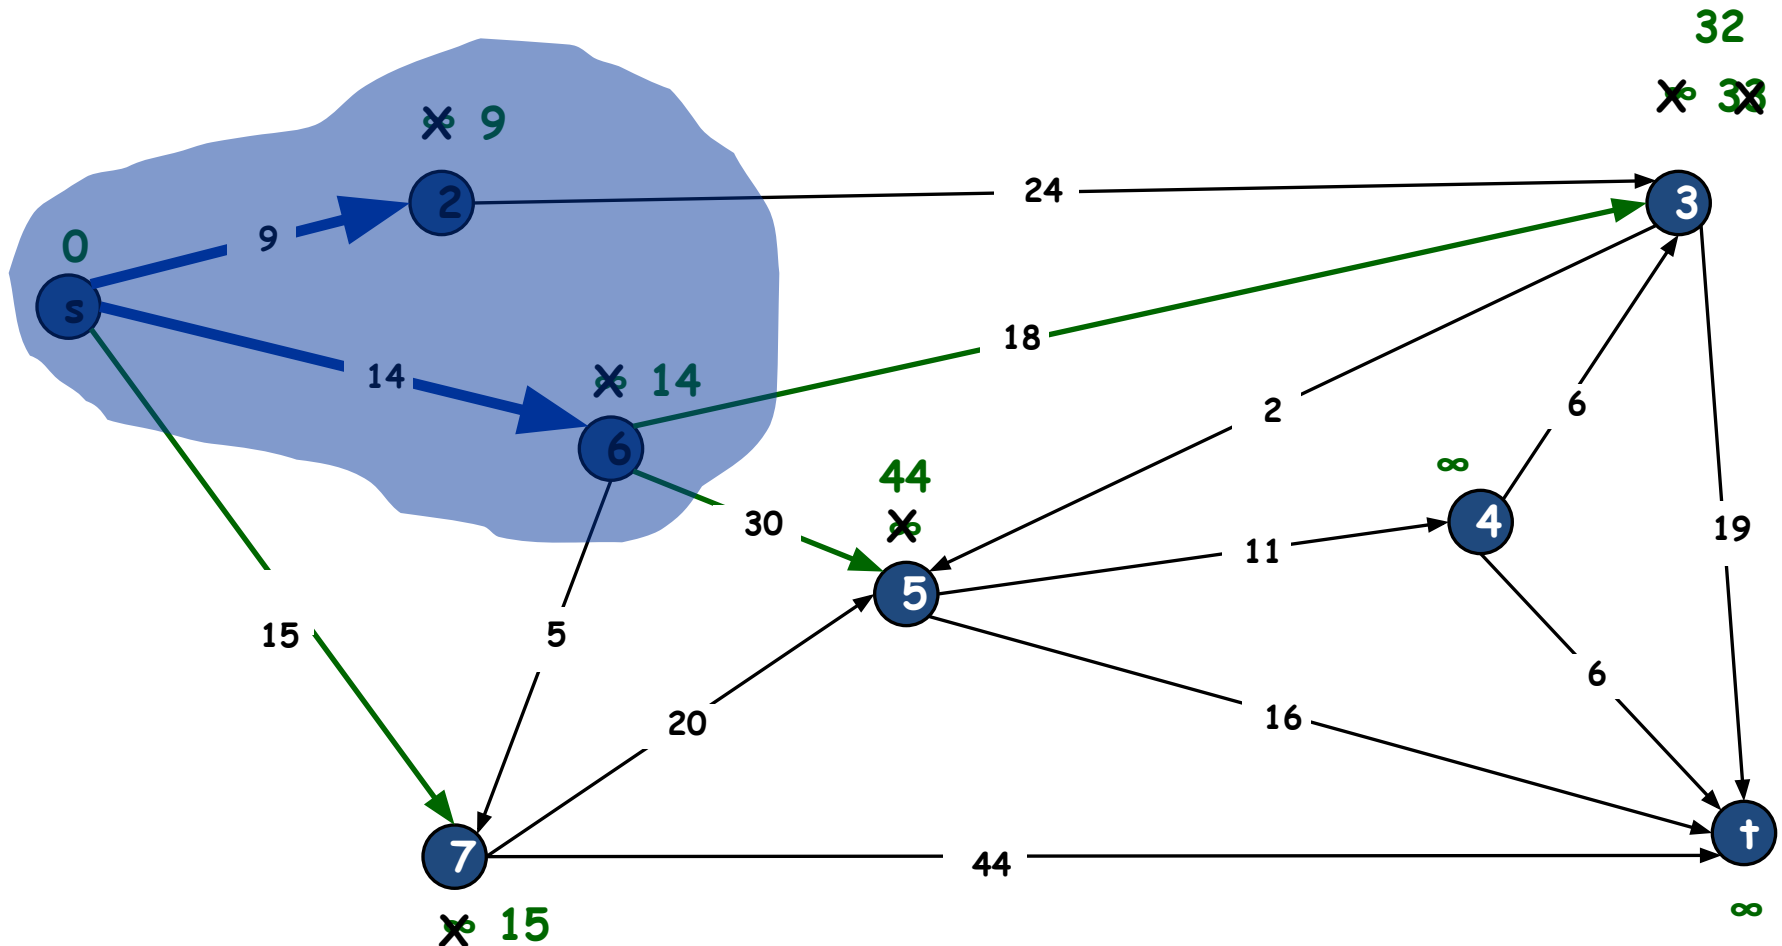
$PQ = \{3, 4, 5, 6, 7, \dagger\}$



Dijkstra's Shortest Path Algorithm

$S = \{s, 2, 6\}$


$PQ = \{3, 4, 5, 7, \dagger\}$



Dijkstra's Shortest Path Algorithm

$S = \{s, 2, 6\}$


$PQ = \{3, 4, 5, 7, \dagger\}$



Dijkstra's Shortest Path Algorithm

$S = \{s, 2, 6, 7\}$


$PQ = \{3, 4, 5, \dagger\}$



Dijkstra's Shortest Path Algorithm

$S = \{s, 2, 6, 7\}$


$PQ = \{3, 4, 5, \dagger\}$



Dijkstra's Shortest Path Algorithm

$S = \{s, 2, 3, 6, 7\}$


$PQ = \{4, 5, \dagger\}$



Dijkstra's Shortest Path Algorithm

$S = \{s, 2, 3, 6, 7\}$

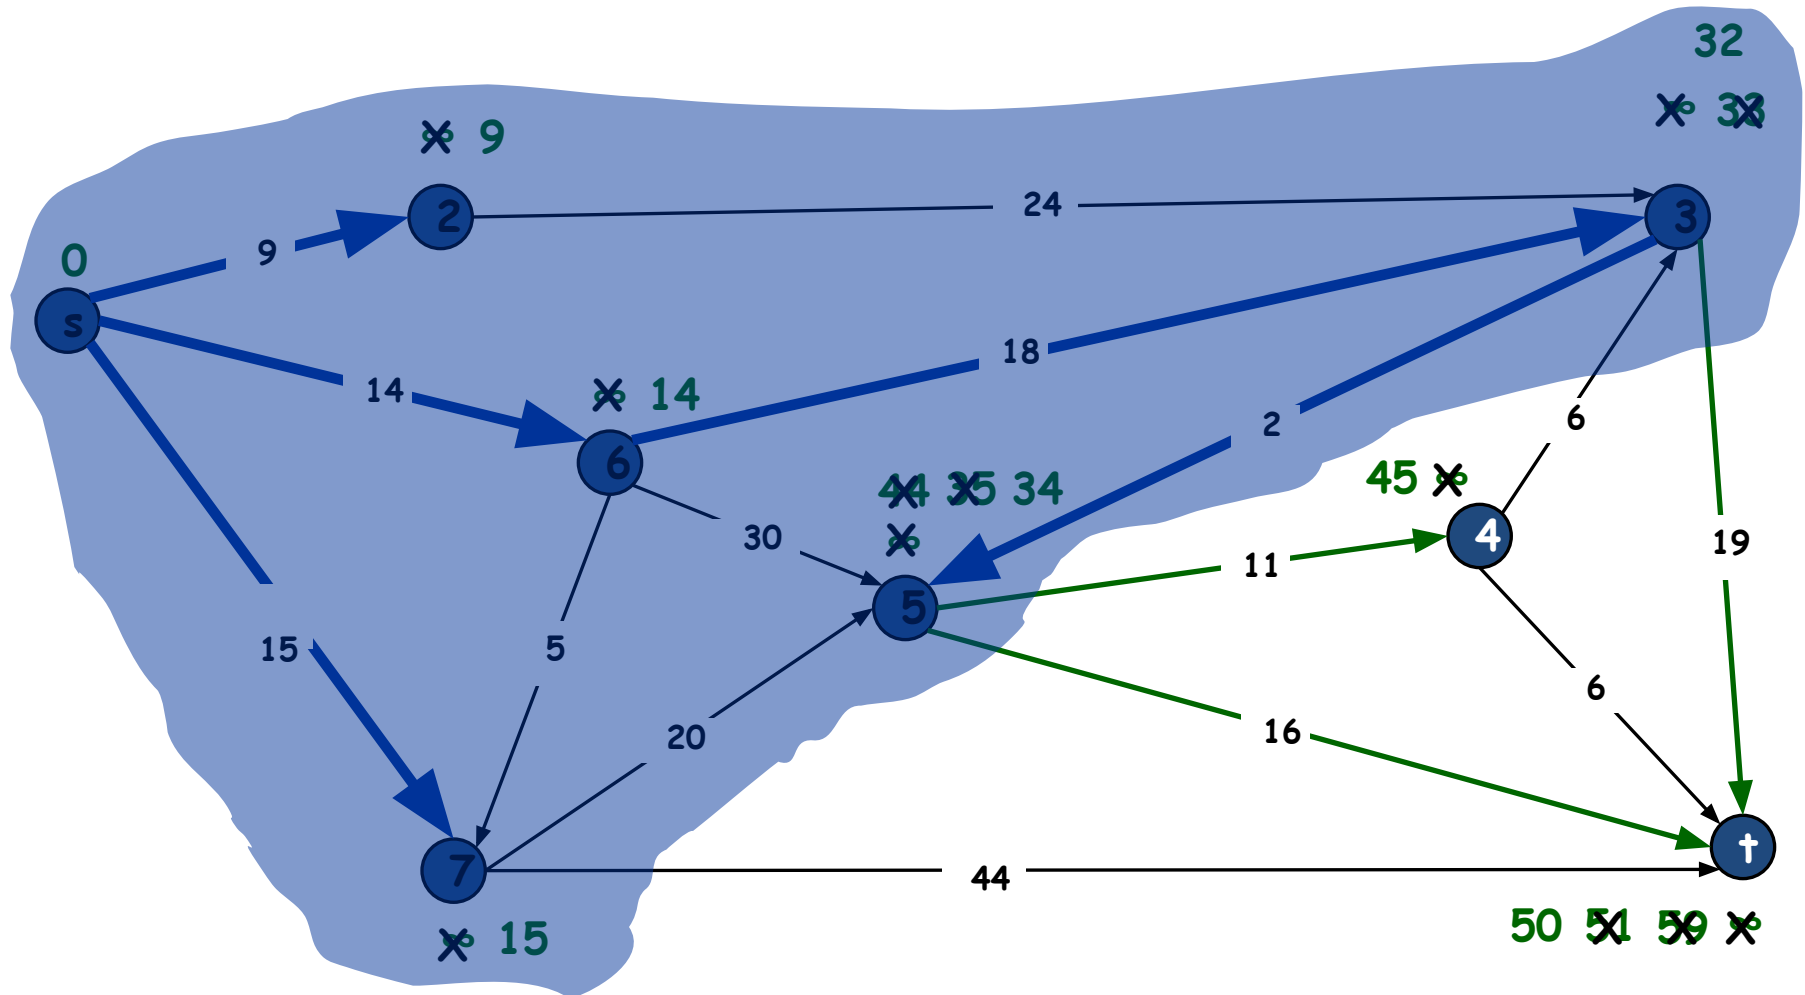
$PQ = \{4, 5, \dagger\}$



Dijkstra's Shortest Path Algorithm

$S = \{s, 2, 3, 5, 6, 7\}$


$PQ = \{4, \dagger\}$



Dijkstra's Shortest Path Algorithm

$S = \{s, 2, 3, 5, 6, 7\}$

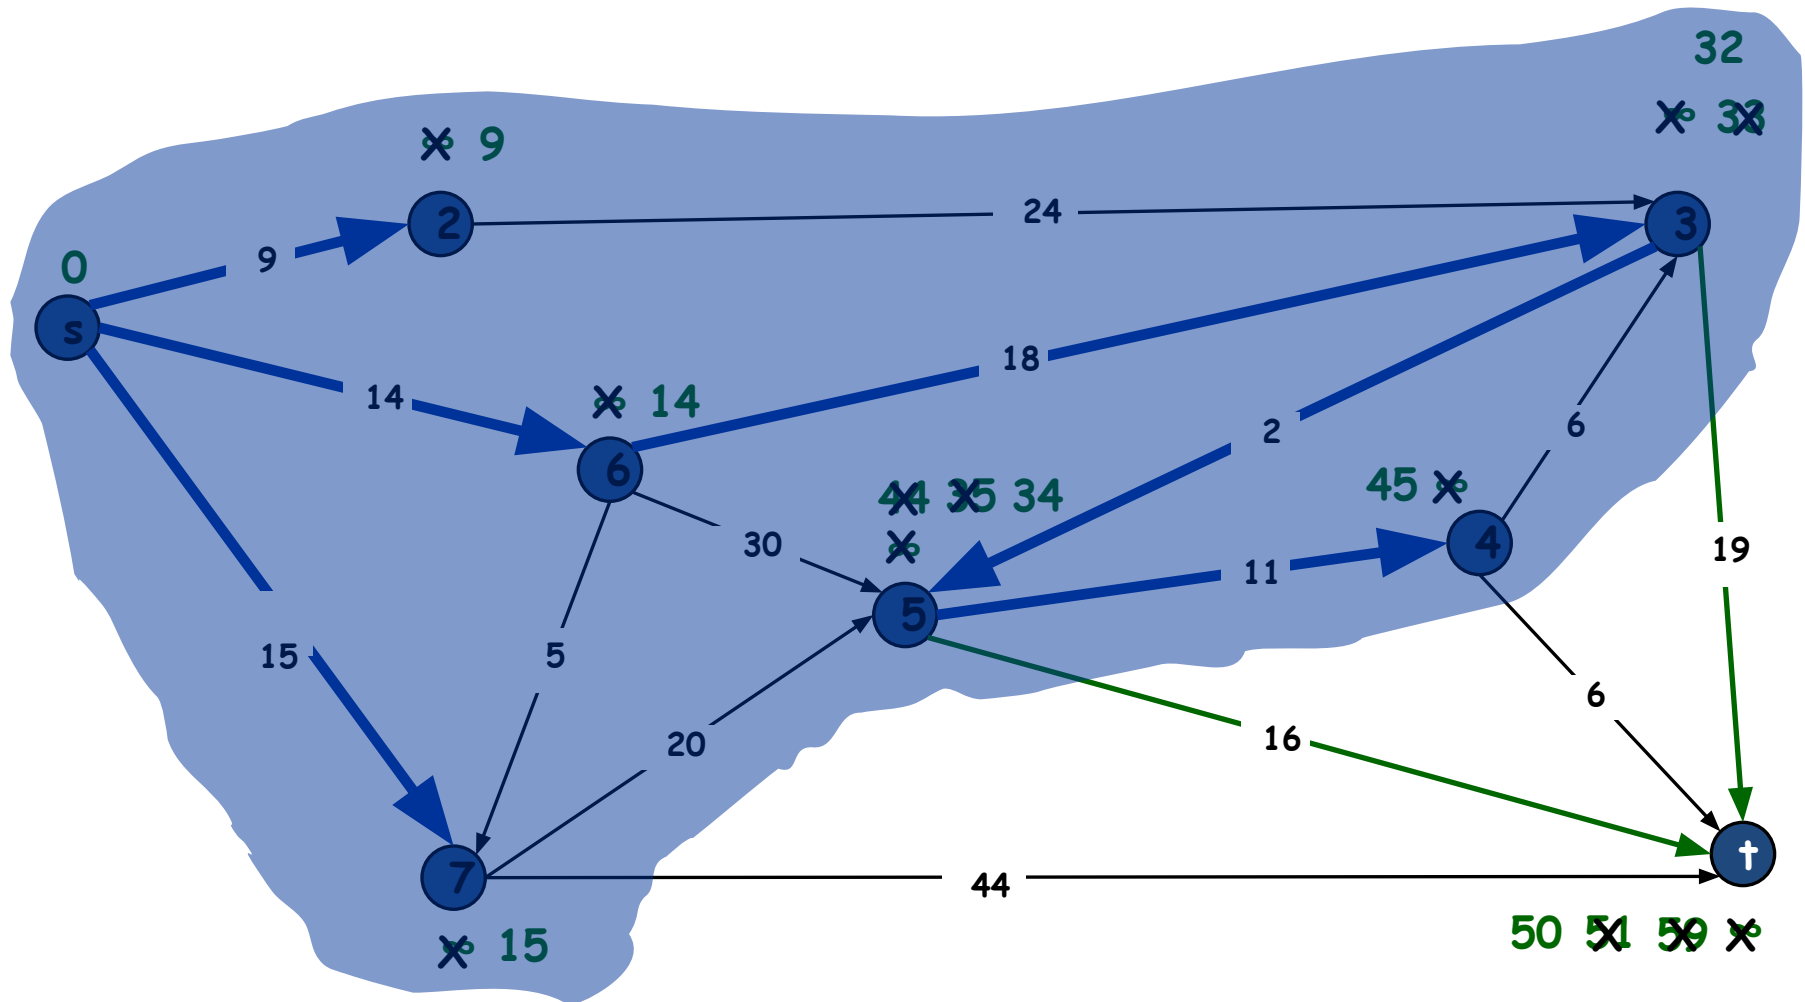
$PQ = \{4, \dagger\}$



Dijkstra's Shortest Path Algorithm

$S = \{s, 2, 3, 4, 5, 6, 7\}$

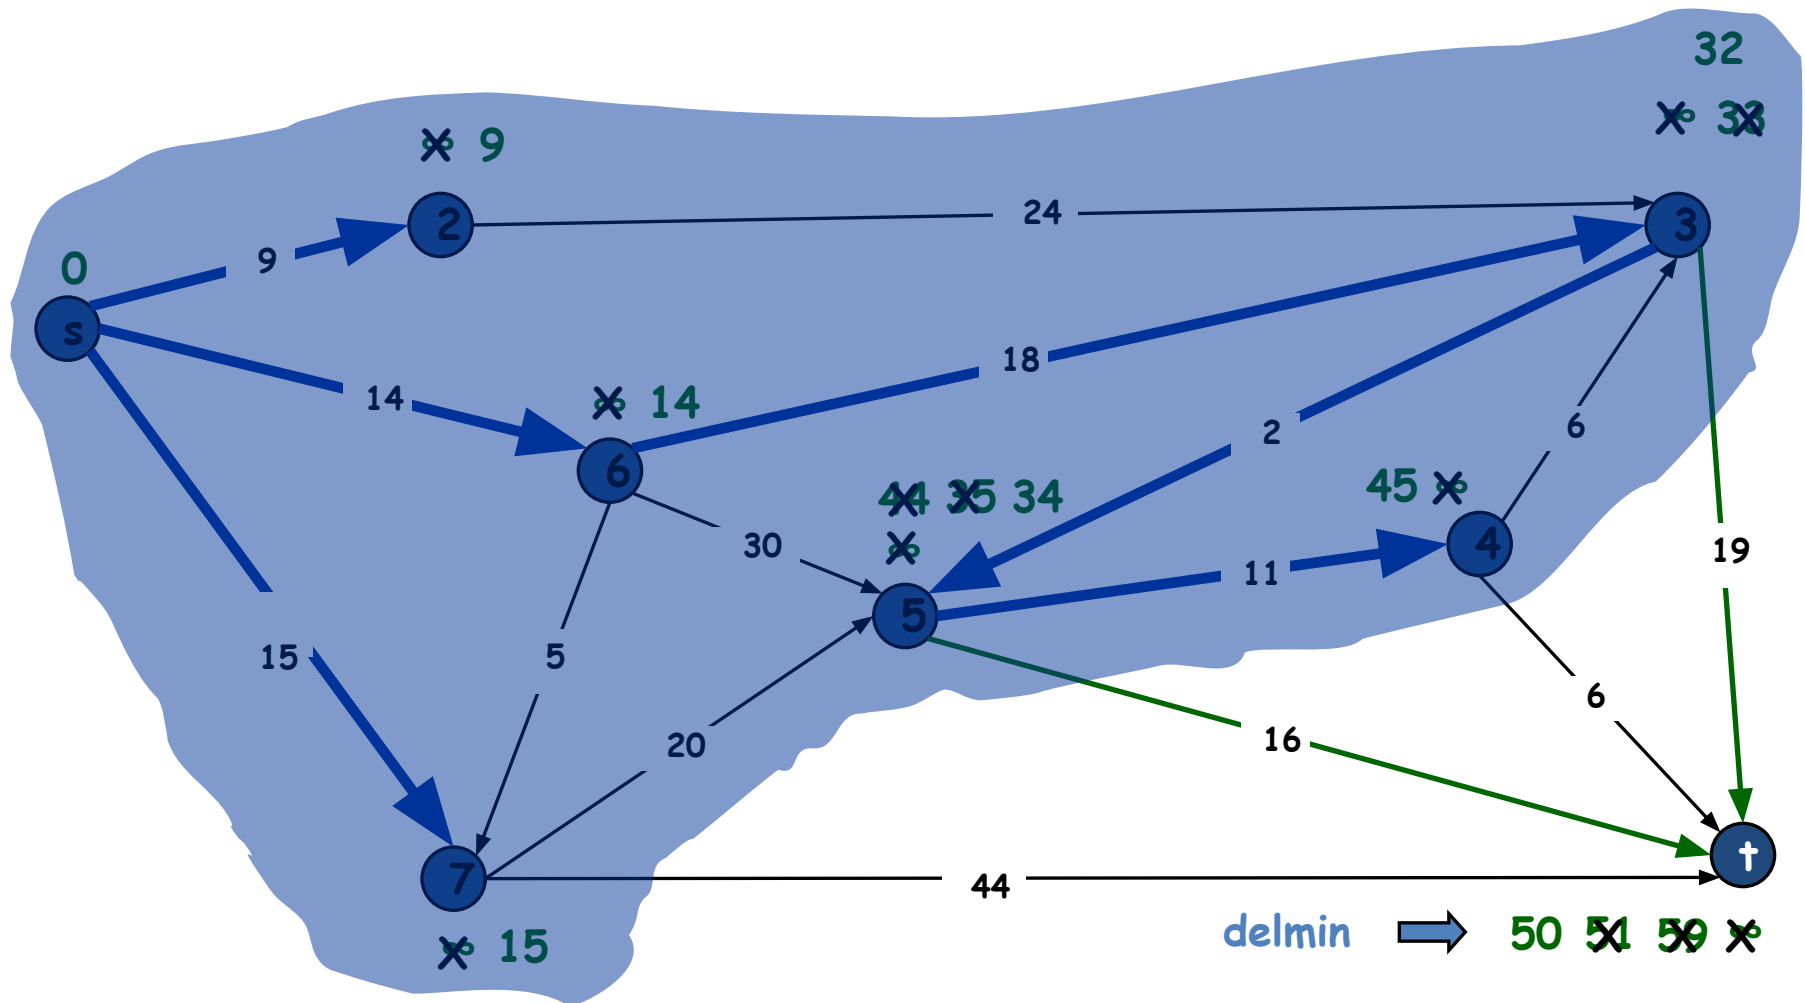
$PQ = \{t\}$



Dijkstra's Shortest Path Algorithm

$S = \{s, 2, 3, 4, 5, 6, 7\}$

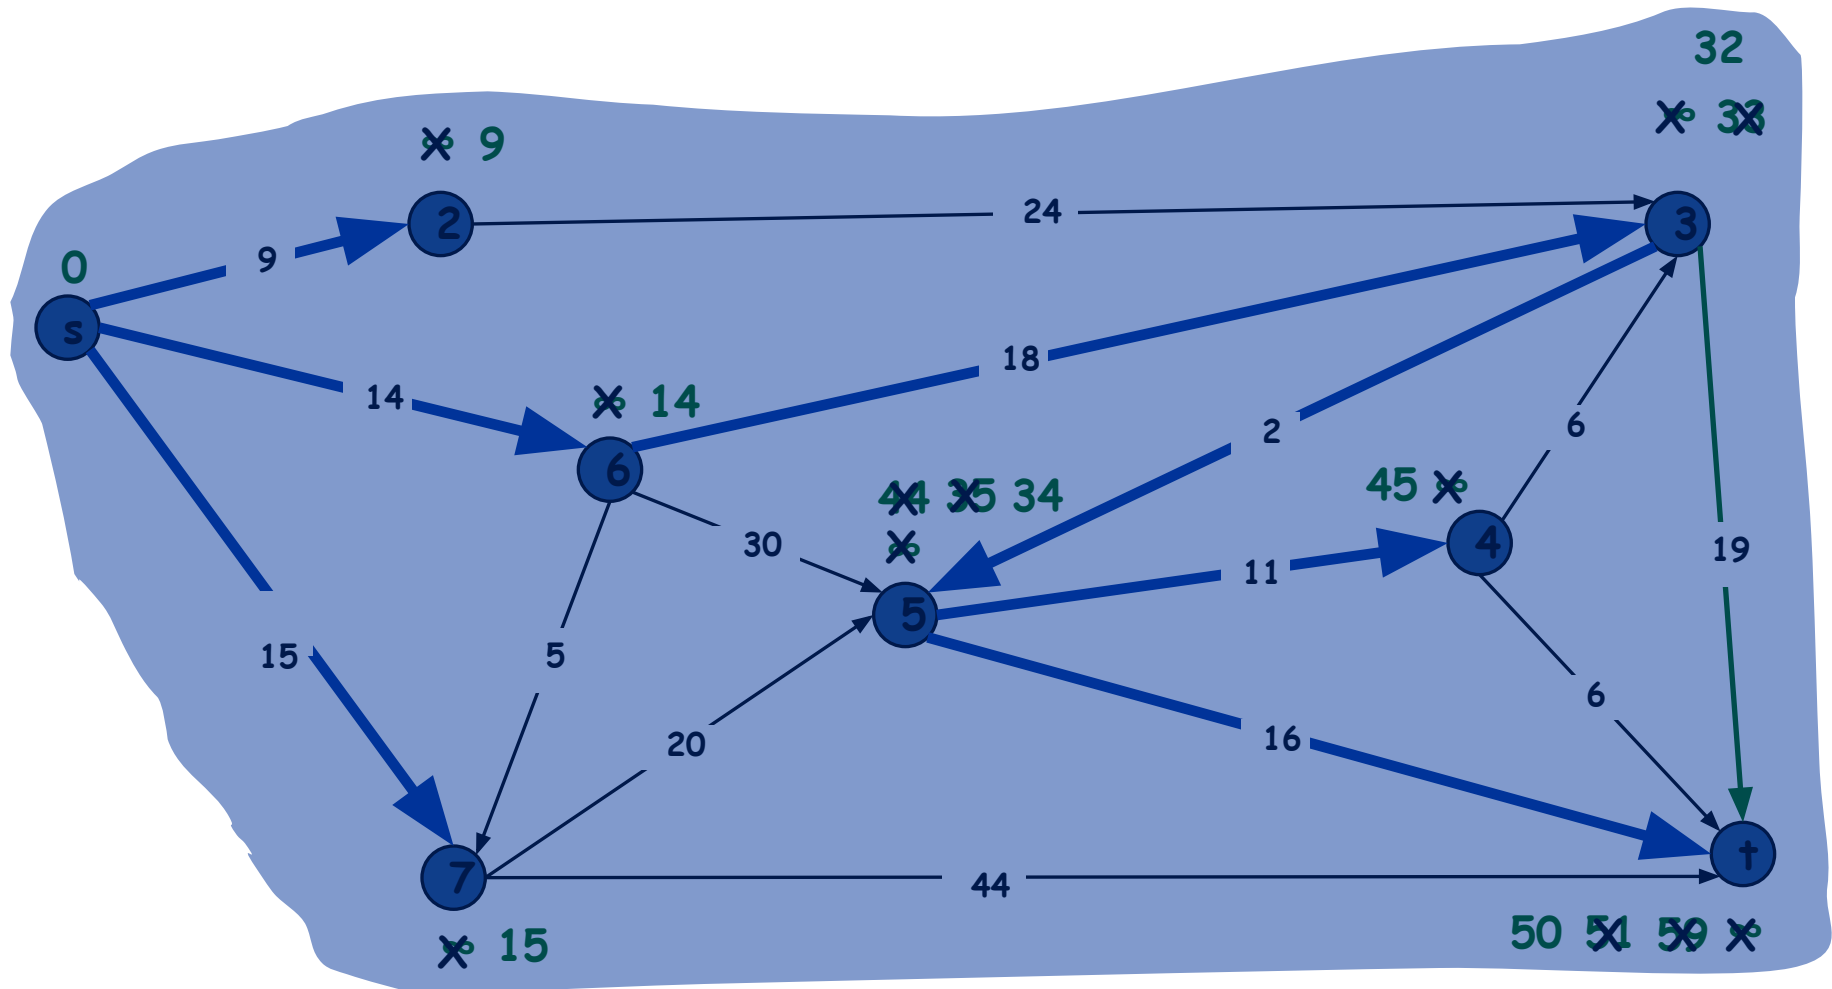
$PQ = \{t\}$



Dijkstra's Shortest Path Algorithm

$S = \{s, 2, 3, 4, 5, 6, 7, t\}$

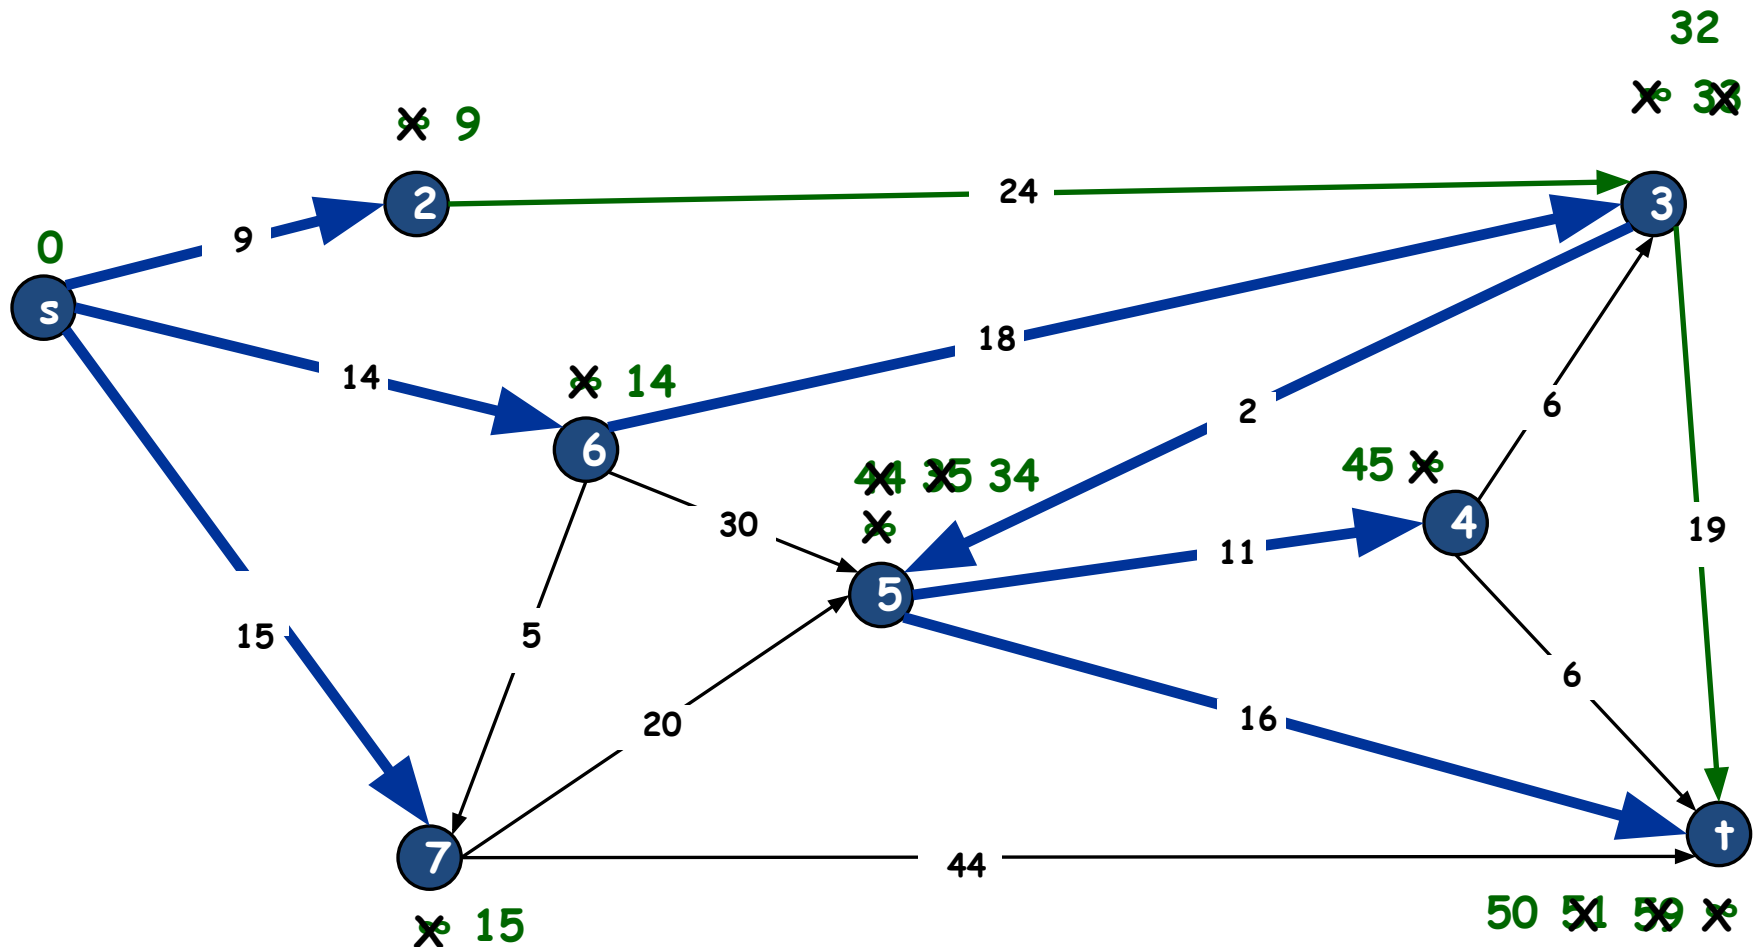
$PQ = \{\}$



Dijkstra's Shortest Path Algorithm

$S = \{s, 2, 3, 4, 5, 6, 7, t\}$

$PQ = \{\}$



Dijkstra's algorithm: efficient implementation

- **Implementation.**
 - Algorithm maintains $\pi[v]$ for each node v .
 - Priority queue stores unexplored nodes, using $\pi[\cdot]$ as priorities.
 - Once u is deleted from the PQ, $\pi[u]$ = length of a shortest $s \rightsquigarrow u$ path.

DIJKSTRA (V, E, ℓ, s)

FOREACH $v \neq s : \pi[v] \leftarrow \infty, \text{pred}[v] \leftarrow \text{null}; \pi[s] \leftarrow 0.$

Create an empty priority queue PQ .

FOREACH $v \in V : \text{INSERT}(PQ, v, \pi[v]).$

WHILE (*IS-NOT-EMPTY*(PQ))

$u \leftarrow \text{DEL-MIN}(PQ).$

FOREACH edge $e = (u, v) \in E$ leaving u :

IF ($\pi[v] > \pi[u] + \ell_e$)

$\pi[v] \leftarrow \pi[u] + \ell_e; \text{pred}[v] \leftarrow e.$

DECREASE-KEY($PQ, v, \pi[v]$).

Dijkstra's algorithm: which priority queue?

- **Performance.** Depends on PQ: n INSERT, n DELETE-MIN, $\leq m$ DECREASE-KEY.
 - Array implementation optimal for dense graphs. $\leftarrow \Theta(n^2)$ edges
 - Binary heap much faster for sparse graphs. $\leftarrow \Theta(n)$ edges
 - 4-way heap worth the trouble in performance-critical situations.


priority queue	INSERT	DELETE-MIN	DECREASE-KEY	total
unordered array	$O(1)$	$O(n)$	$O(1)$	$O(n^2)$
binary heap	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(m \log n)$
d-way heap (Johnson 1975)	$O(d \log n)$	$O(d \log_d n)$	$O(\log_d n)$	$O(m \log_{m/n} n)$
Fibonacci heap (Fredman–Tarjan 1984)	$O(1)$	$O(\log n)^\dagger$	$O(1)^\dagger$	$O(m + n \log n)$
integer priority queue (Thorup 2004)	$O(1)$	$O(\log \log n)$	$O(1)$	$O(m + n \log \log n)$

MINIMUM SPANNING TREES

[KT 4.5]

Adapted from Slides by
Kevin Wayne.
Copyright © 2005 Pearson-Addison Wesley.
All rights reserved.

And Bistra Dilkina, Anne Benoit



Trees


Def. A **path** is a sequence edges which connects a sequence of nodes.

Def. A **cycle** is a path with no repeated nodes or edges other than the start and end nodes.

Def. An undirected graph is a **tree** if it is connected and does not contain a cycle.

Theorem. Let G be an undirected graph on n nodes. Any two of the following statements imply the third.


- G is connected.
- G does not contain a cycle.
- G has $n-1$ edges.



Spanning Tree

Def. An undirected graph is a **spanning tree** if it is a tree and touches every vertex in G .


Cayley's Theorem. There are n^{n-2} spanning trees of K_n .




$$G = (V, E)$$

Minimum Spanning Tree

Minimum spanning tree. Given a connected graph $G = (V, E)$ with real-valued edge weights c_e , an MST is a subset of the edges $T \subseteq E$ such that T is a spanning tree whose sum of edge weights is minimized.



$G = (V, E)$



$T, \sum_{e \in T} c_e = 50$

Can't solve by brute force

Applications

MST is a fundamental problem with diverse applications.

- Network design.
 - telephone, electrical, TV cable, computer, road
- Approximation algorithms for NP-hard problems.
 - traveling salesperson problem, Steiner tree
- Computational biology (single cells)
 - Eg. reconstruct trajectory of stem cell differentiation
- Cluster analysis.

Greedy Algorithms

Greedy Algorithms

Prim's algorithm. Start with some root node s and greedily grow a tree T from s outward. At each step, add the cheapest edge e to T that has exactly one endpoint in T .

Greedy Algorithms

Prim's algorithm. Start with some root node s and greedily grow a tree T from s outward. At each step, add the cheapest edge e to T that has exactly one endpoint in T .

Kruskal's algorithm. Start with $T = \emptyset$. Consider edges in ascending order of cost. Insert edge e in T unless doing so would create a cycle.

Greedy Algorithms

Prim's algorithm. Start with some root node s and greedily grow a tree T from s outward. At each step, add the cheapest edge e to T that has exactly one endpoint in T .

Kruskal's algorithm. Start with $T = \emptyset$. Consider edges in ascending order of cost. Insert edge e in T unless doing so would create a cycle.

Reverse-Delete algorithm. Start with $T = E$. Consider edges in descending order of cost. Delete edge e from T unless doing so would disconnect T .

Greedy Algorithms

Prim's algorithm. Start with some root node s and greedily grow a tree T from s outward. At each step, add the cheapest edge e to T that has exactly one endpoint in T .

Kruskal's algorithm. Start with $T = \emptyset$. Consider edges in ascending order of cost. Insert edge e in T unless doing so would create a cycle.


Reverse-Delete algorithm. Start with $T = E$. Consider edges in descending order of cost. Delete edge e from T unless doing so would disconnect T .

Remark. All three algorithms produce an MST.

Greedy Algorithms


Def. A **cut** is a partition of the nodes into two nonempty subsets S and $V - S$.

Def. The **cutset** of a cut S is the set of edges with exactly one endpoint in S .



Cut $S = \{4, 5, 8\}$

Cutset $D = 5-6, 5-7, 3-4, 3-5, 7-8$



e is in the MST


Greedy Algorithms

Def. A **cut** is a partition of the nodes into two nonempty subsets S and $V - S$.

Def. The **cutset** of a cut S is the set of edges with exactly one endpoint in S .


Simplifying assumption. All edge costs c_e are distinct.

Cut property. Let S be any subset of nodes, and let e be the min cost edge with exactly one endpoint in S . Then every MST contains e .



Cut $S = \{4, 5, 8\}$

Cutset $D = \{5-6, 5-7, 3-4, 3-5, 7-8\}$



e is in the MST

Proof of the cut property


Cut Property

Cut property. Let S be **any** subset of nodes, and let $e = (u, v)$ be the min cost edge with exactly one endpoint in S . Then every MST T^* contains e .

Simplifying assumption. All edge costs c_e are distinct.

Pf. (exchange argument)


- Given: e is the min cost edge in the cutset for a set S , T^* is an MST
- Suppose e does not belong to T^* , and let's see what happens.
- The endpoints u and v of e must be connected by a path in T^* (spanning)
- \Rightarrow Adding e to T^* creates a cycle C in T^*
- Edge e is both in the cycle C and in the cutset corresponding to $S \Rightarrow$ there exists another edge, say f , **that is in both C and cutset of S** .
- $T' = T^* \cup \{e\} - \{f\}$ is also a spanning tree, because:
 - All n nodes are connected (f, e were part of a cycle) AND $n - 1$ edges
- Since $c_e < c_f$, $\text{cost}(T') < \text{cost}(T^*)$.
- This is a contradiction since T^* is an optimal spanning tree, our assumption about e not being in T^* must be wrong.



Prim's Algorithm

Prim's algorithm. [Jarník 1930, Dijkstra 1957, Prim 1959]

- Initialize S = any node.
- (Apply cut property to S .)
- Add to tree the **min cost edge (u, v) in cutset** corresponding to S , and add one new explored node u to S .




Prim's Algorithm: Proof of Correctness

Prim's algorithm. [Jarník 1930, Dijkstra 1957, Prim 1959]

- Initialize $S = \text{any node}$.
- (Apply cut property to S .)
- Add to tree the min cost edge (u, v) in cutset corresponding to S , and add one new explored node u to S .

1) Every edge added satisfies the Cut Property: by design

2) It produces a spanning tree: alg stops when $S = V$



Implementation: Prim's Algorithm

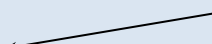
Implementation: Prim's Algorithm

Implementation. Use a priority queue (as for Dijkstra).

- Maintain set of explored nodes S .
- For each unexplored node v , maintain attachment cost $a[v] = \text{cost of cheapest edge } v \text{ to a node in } S$.
- $O(n^2)$ with an array; $O(m \log n)$ with a binary heap.

```
Prim(G, c) {  
    foreach ( $v \in V$ )  $a[v] \leftarrow \infty$   
    Initialize an empty priority queue  $Q$   
    foreach ( $v \in V$ ) insert  $v$  onto  $Q$   
    Initialize set of explored nodes  $S \leftarrow \emptyset$   
  
    while ( $Q$  is not empty) {  
         $u \leftarrow$  delete min element from  $Q$   
         $S \leftarrow S \cup \{u\}$   
        foreach (edge  $e = (u, v)$  incident to  $u$ )  
            if ( $(v \notin S)$  and  $(c_e < a[v])$ )  
                decrease priority  $a[v]$  to  $c_e$   
    }  
}
```

Cheapest edge between v and a
node in explored set S




MST: cycle property


Simplifying assumption: All edge costs c_e are distinct.

Cycle. Set of edges the form $a-b, b-c, c-d, \dots, y-z, z-a$.

Cycle property. Let C be any cycle, and let f be the max cost edge belonging to C . Then every MST does not contain f .



Cycle $C = 1-2, 2-3, 3-4, 4-5, 5-6, 6-1$



f is not in the MST


Cycle property: proof

Cycle property. Let C be any cycle in G , and let f be the max cost edge belonging to C . Then every MST T^* does not contain f .

Simplifying assumption. All edge costs c_e are distinct.

Pf. (exchange argument)

- Given: f is the max cost edge in a cycle C , T^* is an MST
- Suppose f belongs to T^* , and let's see what happens.
- Deleting f from T^* disconnects T^* and creates a cut S in T^* .
- Edge f is both in the cycle C and in the cutset D corresponding to S
 \Rightarrow there exists another edge, say e , that is in both C and D .
- $T' = T^* \cup \{e\} - \{f\}$ is also a spanning tree (same argument as before).
- Since $c_e < c_f$, $\text{cost}(T') < \text{cost}(T^*)$.
- This is a contradiction since T^* is an optimal spanning tree, hence our assumption about f is wrong.



Reverse-Delete Algorithm: Proof of Correctness

The algorithm. Start with $T = E$. Consider edges in descending order of cost. Delete edge e from T unless doing so would disconnect T .

Th. This algorithm produces an MST of G .

Cycle property. Let C be any cycle in G , and let f be the max cost edge belonging to C . Then every MST T^* does not contain f .

Pf.

- Edges removed do not belong to any MST.
- The output is a spanning tree:
connected & no cycle.

